

Predict Taxi Fares with Random Forests

Introduction

To drive a yellow New York taxi, you have to hold a "medallion" from the city's *Taxi and Limousine Commission*. Recently, one of those changed hands for over one million dollars, which shows how lucrative the job can be. But this is the age of business intelligence and analytics! Even taxi drivers can stand to benefit from some careful investigation of the data, guiding them to maximize their profits. We will also use regression trees and random forests to build a model to predict the locations and times when the biggest fares can be earned.

The dataset

In this project, we will analyze a random sample of 49999 New York journeys made in 2013. It has the following columns:

- medallion: Medallion ID.
- pickup_datetime: The time when the taxi driver pick up the customer.
- pickup_longitude: The longitude of the location where the taxi driver pick up the customer.
- pickup_latitude: The latitude of the location where the taxi driver pick up the customer.
- Trip_time_in_secs: Customer trip duration in second.
- fare_amount: The amount of fare paid by the customer.
- tip_amount: The amount of tip paid by the customer.

Let's load the dataset and packages, and take a look at the first few rows.

```
```{r}
library(tidyverse)
library(ggmap)
library(viridis)
library(tree)
library(lubridate)
library(randomForest)
library(party)
taxi <- read_csv("taxi.csv")
head(taxi)
```
```

| | medallion | pickup_datetime | pickup_longitude | pickup_latitude | trip_time_in_secs | fare_amount | tip_amount |
|----|----------------------------------|---------------------|------------------|-----------------|-------------------|-------------|------------|
| 1 | 4D24F4D8EF35878595044A52B098DFD2 | 2013-01-13 10:23:00 | -73.94646 | 40.77273 | 600 | 8.0 | 2.50 |
| 2 | A49C37EB966E7B05E69523D1CB7BE303 | 2013-01-13 04:52:00 | -73.99827 | 40.74041 | 840 | 18.0 | 0.00 |
| 3 | 1E4B72A8E623888F53A9693C364AC05A | 2013-01-13 10:47:00 | -73.95346 | 40.77586 | 60 | 3.5 | 0.70 |
| 4 | F7E4E9439C46B8AD5B16AB9F1B3279D7 | 2013-01-13 11:14:00 | -73.98137 | 40.72473 | 720 | 11.5 | 2.30 |
| 5 | A9DC75D59E0EA27E1ED328E8BE8CD828 | 2013-01-13 11:24:00 | -73.96800 | 40.76000 | 240 | 6.5 | 0.00 |
| 6 | 19BF1BB516C4E992EA3FBAEDA73D6262 | 2013-01-13 10:51:00 | -73.98502 | 40.76341 | 540 | 8.5 | 1.70 |
| 7 | 5F2EFC03B544635C9B0E7A4AA4FF9AC3 | 2013-01-13 12:53:00 | -73.97295 | 40.79527 | 0 | 2.5 | 0.00 |
| 8 | 8DEB70907D00AA1D7FF5E2683240549B | 2013-01-13 07:59:00 | -73.96577 | 40.76530 | 120 | 4.0 | 0.00 |
| 9 | E15F7CC808DD15E0496D830D3DEDECE | 2013-01-13 08:09:00 | -73.94768 | 40.77507 | 720 | 14.0 | 2.00 |
| 10 | 0B3D3D51C78E944F68DC04209E86D5F7 | 2013-01-13 12:53:00 | -73.98457 | 40.72488 | 180 | 4.0 | 3.00 |
| 11 | 3C16CFAD2B12F3508F7211C37F8F888F | 2013-01-13 13:07:00 | -73.97068 | 40.78506 | 360 | 6.5 | 0.00 |

As you can see above, the taxi dataset contains the times and prices of a large number of taxi trips. We also got to know the location, the longitude and latitude, where the trip was started.

Cleaning data is a large part of any data scientist's daily work. It may not seem glamorous, but it makes the difference between a successful model and a failure. The taxi dataset needs a bit of polishing before we're ready to use it.

```
```{r}
Renaming the location variables,
dropping any journeys with zero fares and zero tips,
and creating the total variable as the log sum of fare and tip
taxi <- taxi %>%
 rename(long = pickup_longitude, lat = pickup_latitude) %>%
 filter(fare_amount>0 | tip_amount>0) %>%
 mutate(total = log(fare_amount + tip_amount)) # We use log to reduce the skew
```
```

| | medallion | pickup_datetime | long | lat | trip_time_in_secs | fare_amount | tip_amount | total |
|----|-----------------------------------|---------------------|-----------|----------|-------------------|-------------|------------|-----------|
| 1 | 4D24F4D8EF35878595044A52B098DFD2 | 2013-01-13 10:23:00 | -73.94646 | 40.77273 | 600 | 8.0 | 2.50 | 2.3513753 |
| 2 | A49C37EB966E7B05E69523D1CB7BE303 | 2013-01-13 04:52:00 | -73.99827 | 40.74041 | 840 | 18.0 | 0.00 | 2.8903718 |
| 3 | 1E4B72A8E623888F53A9693C364AC05A | 2013-01-13 10:47:00 | -73.95346 | 40.77586 | 60 | 3.5 | 0.70 | 1.4350845 |
| 4 | F7E4E9439C46B8AD5B16AB9F1B3279D7 | 2013-01-13 11:14:00 | -73.98137 | 40.72473 | 720 | 11.5 | 2.30 | 2.6246686 |
| 5 | A9DC75D59E0EA27E1ED328E8BE8CD828 | 2013-01-13 11:24:00 | -73.96800 | 40.76000 | 240 | 6.5 | 0.00 | 1.8718022 |
| 6 | 19BF1BB516C4E992EA3FBAEDA73D6262 | 2013-01-13 10:51:00 | -73.98502 | 40.76341 | 540 | 8.5 | 1.70 | 2.3223877 |
| 7 | 5F2EFC03B544635C9B0E7A4AA4FF9AC3 | 2013-01-13 12:53:00 | -73.97295 | 40.79527 | 0 | 2.5 | 0.00 | 0.9162907 |
| 8 | 8DEB70907D00AA1D7FF5E2683240549B | 2013-01-13 07:59:00 | -73.96577 | 40.76530 | 120 | 4.0 | 0.00 | 1.3862944 |
| 9 | E15F7CCB808DD15E0496D830D3DEDECE | 2013-01-13 08:09:00 | -73.94768 | 40.77507 | 720 | 14.0 | 2.00 | 2.7725887 |
| 10 | 0B3D3D51C78E944F68DC04209E86D5F7 | 2013-01-13 12:53:00 | -73.98457 | 40.72488 | 180 | 4.0 | 3.00 | 1.9459101 |
| 11 | 3C16CFAD2B12F3508F7211C37F8F8B8F | 2013-01-13 13:07:00 | -73.97068 | 40.78506 | 360 | 6.5 | 0.00 | 1.8718022 |
| 12 | 5888835B75CF97CBCE738A58484B12A6D | 2013-01-13 12:31:00 | -74.00164 | 40.74414 | 660 | 9.5 | 1.90 | 2.4336134 |
| 13 | 58A836AD639867DB949723BA2A941734 | 2013-01-13 13:37:00 | -73.98800 | 40.74960 | 420 | 7.0 | 0.00 | 1.9459101 |

While the dataset contains taxi trips from all over New York City, the bulk of the trips are to and from Manhattan, so let's focus only on trips initiated there.

```
```{r}
Reducing the data to taxi trips starting in Manhattan
Manhattan is bounded by the rectangle with
latitude from 40.70 to 40.83 and
longitude from -74.025 to -73.93
taxi <- taxi %>%
 filter(between(lat, 40.70, 40.83), between(long, -74.025, -73.93))
```
```

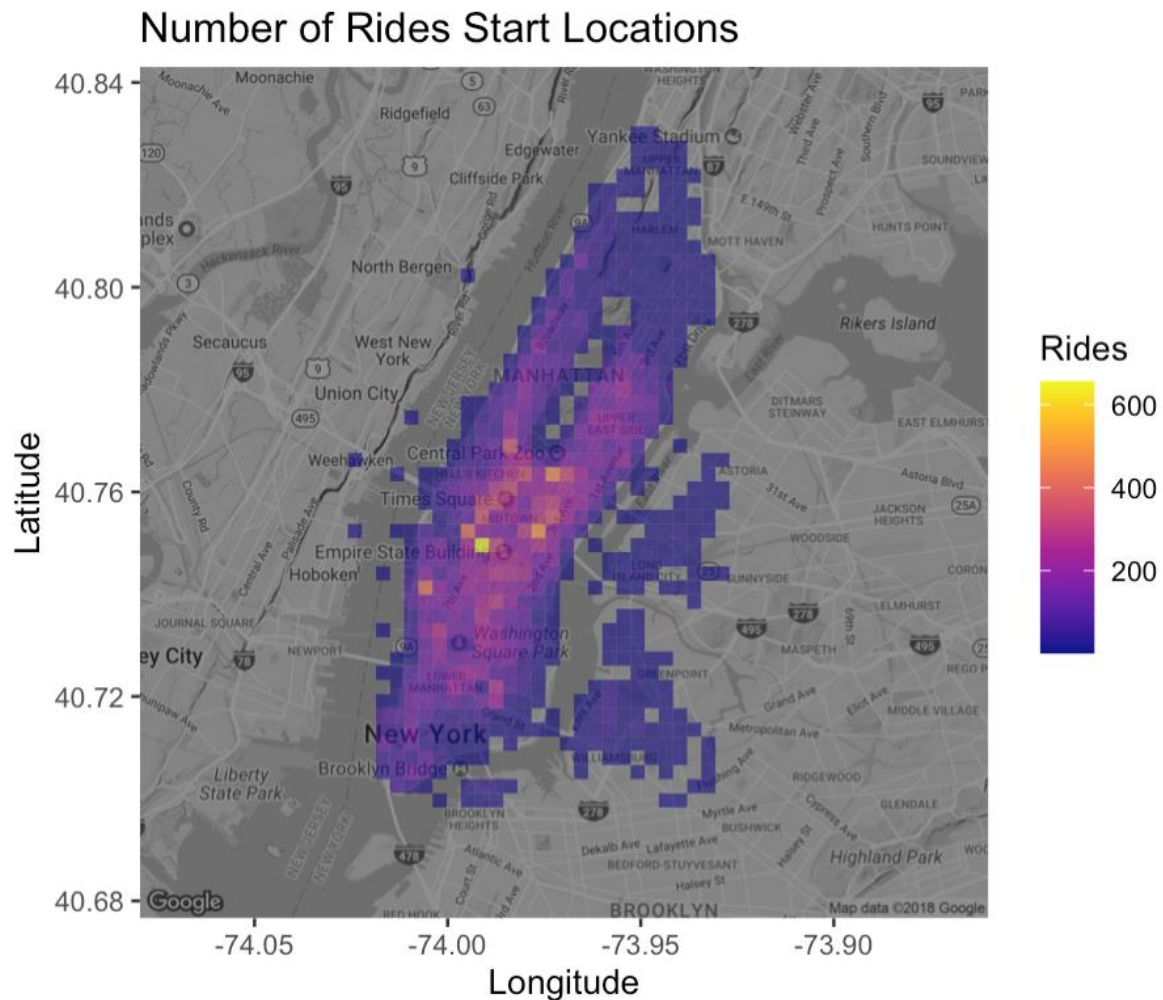
Empirical Analysis

Now we're going to draw a map using ggmap and ggplot2 package to visualize location in Manhattan where people tend to start their taxi trip.

```
```{r}
Drawing a density map with the number of journey start locations
ggmap(manhattan, darken = 0.5) +
 scale_fill_viridis(option = "plasma") +
```

```
geom_bin2d(data = taxi, aes(x = long, y = lat), bins = 60, alpha=0.6) +
labs(x = "Longitude", y = "Latitude", fill = "Rides", title = "Number of Rides Start Locations")
```


```



The map from the previous task showed that the trips are highly concentrated in the business and tourist areas. We also see that some taxi trips originating in Brooklyn slipped through, but that's fine. We're now going to use a regression tree to predict the total fare, with lat and long being the predictors. The tree algorithm will find cutpoints in those predictors that result in the decision tree with the best predictive capability.

```
```{r}
Fitting a tree to lat and long
fitted_tree <- tree(total ~ lat + long, data = taxi)

Draw a diagram of the tree structure
plot(fitted_tree)
text(fitted_tree)
```
```



The tree above looks a bit frugal; it only includes one split: It predicts that trips where $\text{lat} < 40.7237$ are more expensive, which makes sense as it is downtown Manhattan. But that's it. It didn't even include long as tree deemed that it didn't improve the predictions. Taxi drivers will need more information than this, and any driver paying for your data-driven insights would be disappointed. As we know from Robert de Niro, it's best not to upset New York taxi drivers. So we will start by adding some more predictors related to the *time* the taxi trip was made.

```

```{r}
Generate the three new time variables
taxi <- taxi %>%
 mutate(hour = hour(pickup_datetime),
 wday = wday(pickup_datetime, label = TRUE),
 month = month(pickup_datetime, label = TRUE))
```

```

Let's try fitting a new regression tree where we include the new time variables.


```

```{r}
Fitting a tree with total as the outcome and
lat, long, hour, wday, and month as predictors
fitted_tree <- tree(total ~ lat + long + hour + wday + month, data = taxi)

Draw a diagram of the tree structure
plot(fitted_tree)
text(fitted_tree)

Summarizing the performance of the tree
summary(fitted_tree)
```

```



```

Regression tree:
tree(formula = total ~ lat + long + hour + wday + month, data = taxi)
Variables actually used in tree construction:
[1] "lat"
Number of terminal nodes: 2
Residual mean deviance: 0.3041 = 13910 / 45760
Distribution of residuals:
  Min. 1st Qu.  Median    Mean 3rd Qu.   Max.
-1.61900 -0.37880 -0.04244  0.00000  0.32660  2.69900
  
```

The regression tree has not changed after including the three-time variables. This is likely because latitude is still the most promising first variable to split the data on, and after that split, the other variables are not informative enough to be included. A random forest model, where many different trees are fitted to subsets of the data, may consist of the other variables that make it up in some of the trees.

```

```{r}
Fitting a random forest
fitted_forest <- randomForest(total~lat + long+hour+wday+month, data=taxi, ntree=80,
sampsiz = 10000)

Printing the fitted_forest object
summary(fitted_forest)
```
  
```

```

      Length Class  Mode
call      5 -none- call
type      1 -none- character
predicted 45766 -none- numeric
mse       80 -none- numeric
  
```

```

rsq      80 -none- numeric
oob.times 45766 -none- numeric
importance 5 -none- numeric
importanceSD 0 -none- NULL
localImportance 0 -none- NULL
proximity 0 -none- NULL
ntree     1 -none- numeric
mtry      1 -none- numeric
forest    11 -none- list
coefs     0 -none- NULL
y         45766 -none- numeric
test      0 -none- NULL
inbag     0 -none- NULL
terms     3 terms call

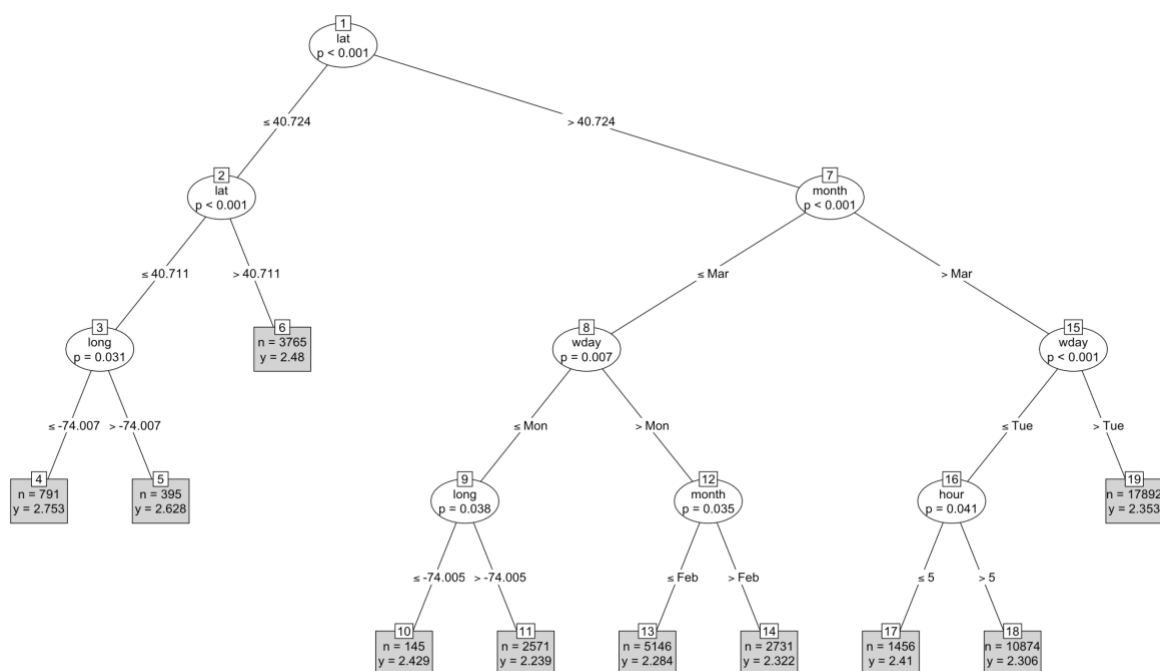
```

Let's draw the random forest plot.

```

```{r}
x <- ctree(total~lat + long+hour+wday+month, data=taxi)
plot(x, type = "simple")
```

```



In the output of `fitted_forest`, you should see the Mean of squared residuals, that is, the average of the squared errors the model makes. If you scroll up and check the summary of `fitted_tree`, you'll find Residual mean deviance, which is the same number. If you compare these numbers, you'll see that `fitted_forest` has a slightly lower error. Neither predictive model is *that* good in statistical terms, they explain only about 3% of the variance.

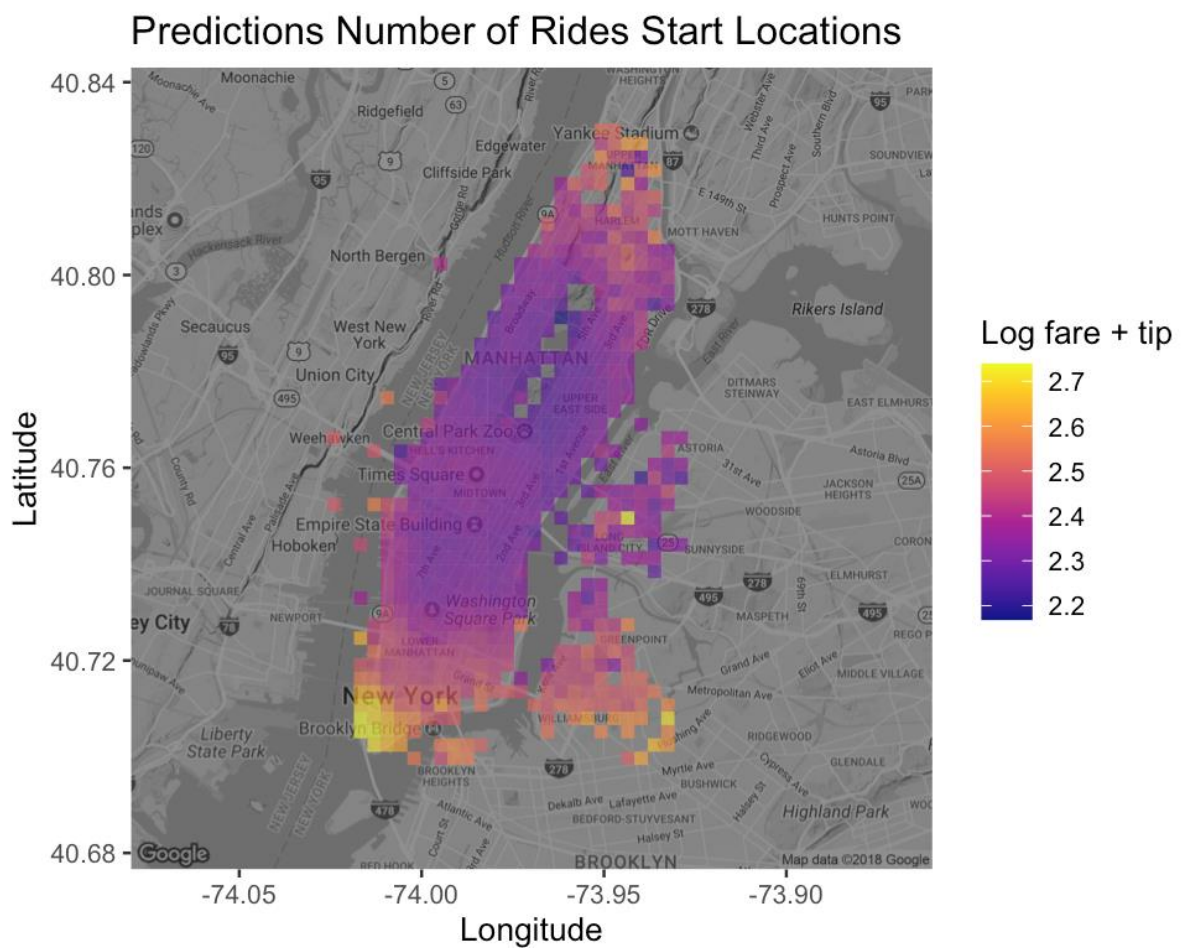
Now, let's take a look at the predictions of `fitted_forest` projected back onto Manhattan.

```

```{r}
Extracting the prediction from fitted_forest
taxi$pred_total <- fitted_forest$predicted

Plotting the predicted mean trip prices from according to the random forest
Drawing a density map with the number of journey start locations
ggmap(manhattan, darken = 0.5) +
 scale_fill_viridis(option = "plasma") +
 stat_summary_2d(data = taxi, aes(x=long, y = lat, z =pred_total), bins = 60, alpha = 0.6, fun
= mean)+
 labs(x="Longitude", y = "Latitude", fill = "Log fare + tip", title = "Predictions Number of
Rides Start Locations ")
```

```



Looking at the map with the predicted fares, we see that taxi fares in downtown Manhattan are predicted to be high, while midtown is lower. This map only shows the prediction as a function of lat and long, but we could also plot the predictions over time or a combination of time and space, but we'll leave that for another time.

For now, let's compare the map with the predicted fares with a new map showing the mean fares according to the data.

```

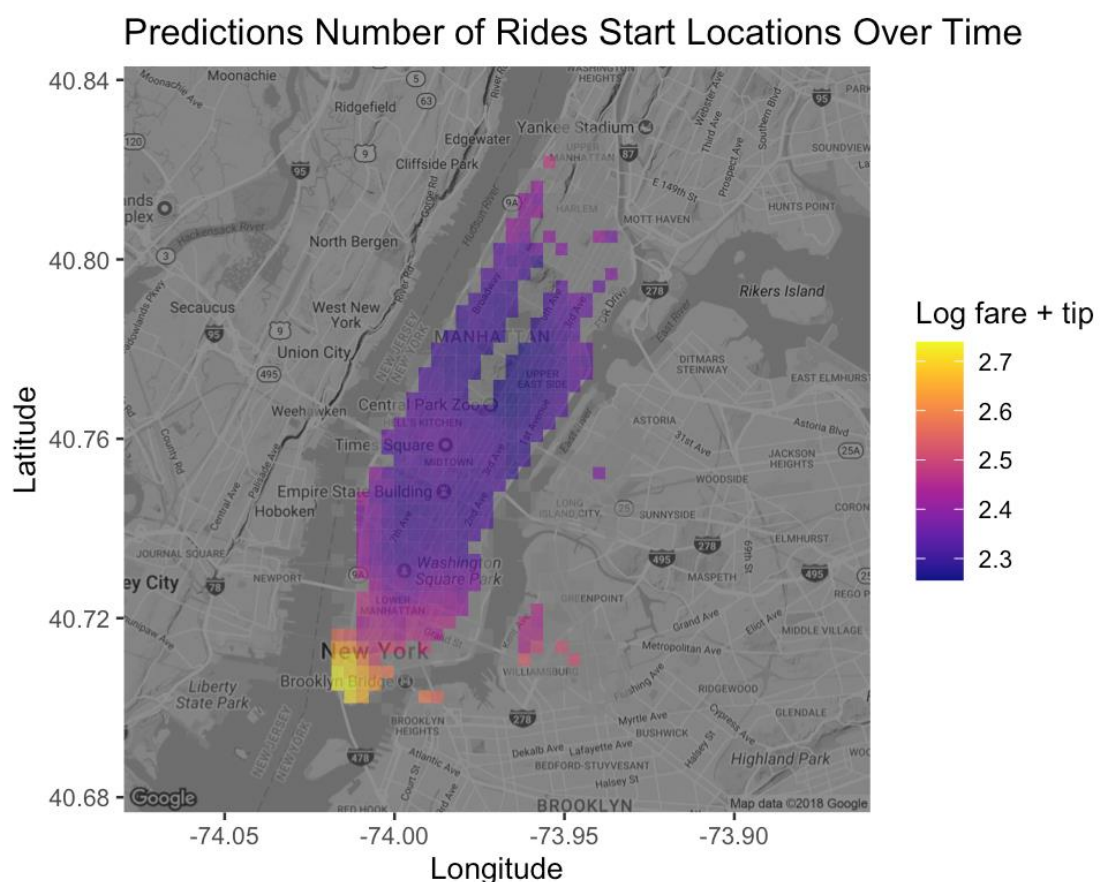
```{r}

```



```
Function that returns the mean *if* there are 15 or more datapoints
mean_if_enough_data <- function(x) {
 ifelse(length(x)>=15, mean(x), NA)
}

Plotting the mean trip prices from the data
ggmap(manhattan, darken = 0.5) +
 scale_fill_viridis(option = "plasma") +
 stat_summary_2d(data = taxi, aes(x=long, y = lat, z =pred_total), bins = 60, alpha = 0.6, fun
= mean_if_enough_data)+
 labs(x="Longitude", y = "Latitude", fill = "Log fare + tip", title = "Predictions Number of
Rides Start Locations Over Time")
` ``
```



## Conclusion

So it looks like the random forest model captured some of the patterns in our data. At this point in the analysis, we could do many more things that we haven't done. We could add more predictors if we have the data. We could try to fine-tune the parameters of randomForest. And we should test the model on a hold-out test dataset. So from the map above, taxi drivers in NYC can expect to spend the most on a taxi ride in lower Manhattan.

## Source

- [https://en.wikipedia.org/wiki/Taxi\\_medallion](https://en.wikipedia.org/wiki/Taxi_medallion)